

How Frappe V14 is bLaZinGLy fast*

* without Rewriting it in Rust



Ankush Menat
Engineer

Frappe

Improvements

- Lower response times
- Lower memory usage
- Lower asset bundle size
- Lower disk usage
- Faster "feels"

How we identify bottlenecks

- Bug reports from paying customers or community
- Frappe Monitor logs using ELK stack
- Sentry
- Developers scratching their itch*

Common patterns

	Knowns	Unknowns
Known	Specific known bottlenecks	Overheads from libraries and abstractions boundaries
Unknown	Our own "small" unknown overheads	<i>Good luck!</i>

Common patterns

	Knowns	Unknowns
Known	Specific known bottlenecks	Overheads from libraries and abstractions boundaries
Unknown	Our own "small" unknown overheads	<i>Good luck!</i>

“Stock Entry submission is slow”

1. Find entry point: **stock_entry.submit()**
2. **%prun sales_invoice.submit()**
3. ...
4. profit?

“Stock Entry submission is slow”

17314183 function calls (17062515 primitive calls) in 18.350 seconds

Ordered by: cumulative time

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
181/1	0.001	0.000	18.350	18.350	{built-in method builtins.exec}
100/1	0.000	0.000	18.350	18.350	document.py:927(submit)
100/1	0.000	0.000	18.350	18.350	document.py:915(_submit)
199/1	0.000	0.000	18.350	18.350	document.py:280
199/1	0.007	0.000	18.350	18.350	document.py:284
2792/15	0.031	0.000	16.988	1.133	document.py:848
2792/15	0.025	0.000	16.809	1.121	document.py:114
2792/15	0.017	0.000	16.791	1.119	document.py:113
501/3	0.002	0.000	16.516	5.505	document.py:854(<lambda>)
299/1	0.007	0.000	12.864	12.864	document.py:984(run_post_save_methods)
1	0.000	0.000	12.519	12.519	stock_entry.py:94(on_submit)

```
if not doctype in self.value_cache:
    self.value_cache = self.value_cache[doctype] = {}
    self.value_cache[doctype] = {}
```

---- clipped to keep relevant output ---

299	0.002	0.000	1.791	0.006	document.py:815(_validate_links)
398	0.042	0.000	1.760	0.004	base_document.py:522(get_invalid_links)

“Stock Entry submission is slow”

☒ Group Duplicate Queries

Index	Query	Duration (ms)	Exact Copies	
2454	SELECT `value` FROM `tabSingles` WHERE `doctype`='Energy Point Settings' AND `field`='enabled'	0.684	301	▼
12	SELECT `name` FROM `tabUOM` WHERE `name` = 'Nos' ORDER BY modified DESC	0.646	299	▼
2163	SELECT `document_type` FROM `tabService Level Agreement` ORDER BY `tabService Level Agreement`.`modified` DESC	0.613	299	▼
2408	SELECT `module`, `custom` FROM `tabDocType` WHERE `name` = 'Stock Ledger Entry' ORDER BY modified DESC	0.625	297	▼
9	SELECT `name` FROM `tabWarehouse` WHERE `name` = 'Stores - _TC' ORDER BY modified DESC	0.662	295	▼
2502	SELECT `module`, `custom` FROM `tabDocType` WHERE `name` = 'Bin' ORDER BY modified DESC	0.576	200	▼
6	SELECT `name` FROM `tabCompany` WHERE `name` = '_Test Company' ORDER BY modified DESC	0.696	199	▼
2425	SELECT `name` FROM `tabDocType` WHERE `name` = 'Stock Entry' ORDER BY modified DESC	0.554	199	▼
2428	SELECT `name` FROM `tabStock Entry` WHERE `name` = 'MAT-STE-2021-00428' ORDER BY modified DESC	0.537	198	▼
2445	SELECT `value` FROM `tabSingles` WHERE `doctype`='Stock Settings' AND `field`='role_allowed_to_create_edit_back_dated_transactions'	0.583	198	▼
2517	SELECT `value` FROM `tabSingles` WHERE `doctype`='Stock Settings' AND `field`='allow_negative_stock'	0.603	198	▼
1769	SELECT field, value FROM `tabSingles` WHERE field in ('valuation_method') AND doctype='Stock Settings'	0.637	197	▼
2442	SELECT `disabled` FROM `tabWarehouse` WHERE `name` = 'Stores - _TC' ORDER BY modified DESC	0.575	196	▼

“Background job crashing from memory usage”

1. Find entry point:

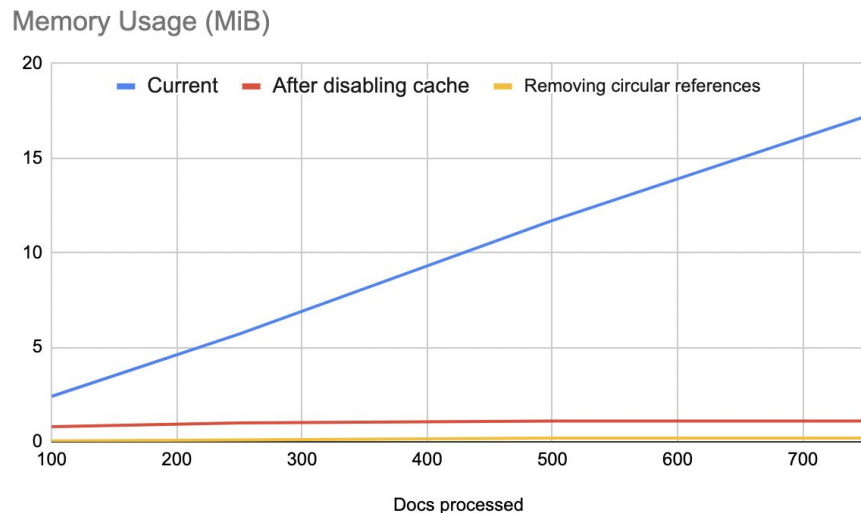
repost_entries()

2. **@memory_profiler.profile**

Line #	Mem usage	Increment	Occurrences	Line Contents
=====				
3	38.816 MiB	38.816 MiB	1	@profile
4				def process_docs():
5	38.828 MiB	0.012 MiB	1	for name in docs_to_process:
6	192.453 MiB	152.625 MiB	1	doc = frappe.get_doc("type", name)
7	204.853 MiB	12.40 MiB	1	doc.process()

“Background job crashing from memory usage”

1. Find entry point:
repost_entries()
2. **@memory_profiler.profile**
3. Keep narrowing down till you find the root cause.
4. ...
5. profit?




Common patterns

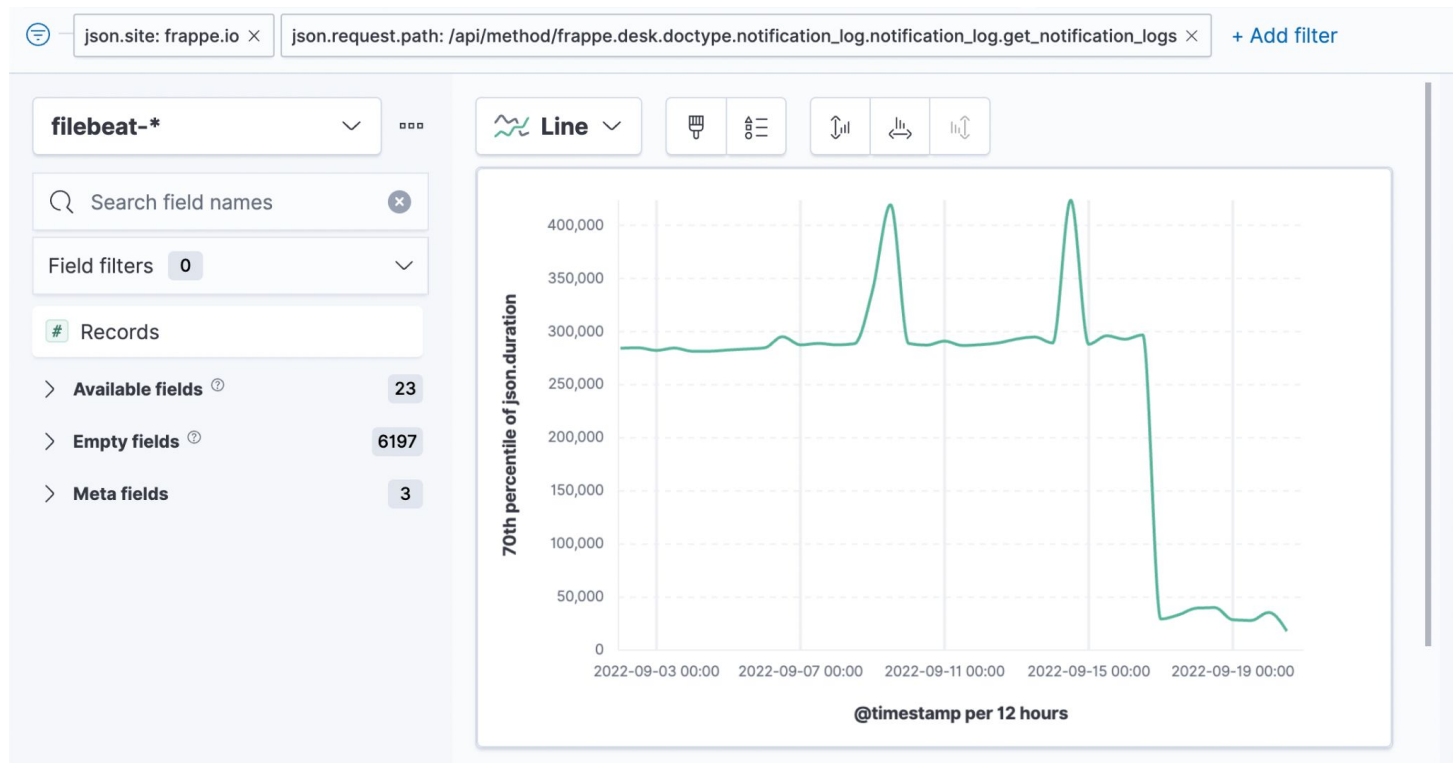
	Knowns	Unknowns
Known	Specific known bottlenecks	Overheads from libraries and abstractions boundaries
Unknown	Our own "small" unknown overheads	<i>Good luck!</i>

“Response times seem slow”

1. Entry point: ??
2. Need high level overview.
3. Use the *right* tool - `py-spy`.
4. ...
5. 30-50% reduction in overheads. profit?

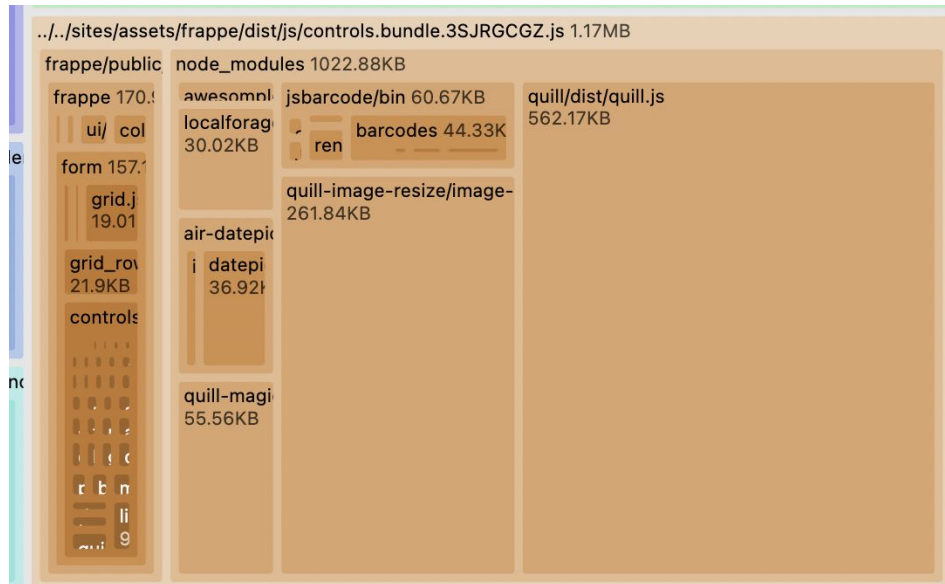


“Response times seem slow”



"Asset bundle size is huge"

1. Use the right tool - ESbuild meta files visualizer
2. Find duplicate, non-critical libraries and remove them
3. ...
4. 5MB -> 3.4 MB. profit?



"Too much disk usage"

1. Sort tables by usage, find top tables
2. Analyze tables.
3. ...
4. 10%-50% drop in db usage..
Profit?

```
{  
  'name' : '29438bf4aa',  
  'ref_doctype' : 'Scheduled Job Log',  
  'docname' : '33588173a2',  
  'data' : '{\n "created_by":  
"Administrator",\n "creation"2022-09-15 02:59:45.896342",\n "updater_reference": null\n}',  
  'doctype' : 'Version'  
}
```

Common patterns

	Knowns	Unknowns
Known	Specific known bottlenecks	Overheads from libraries and abstractions boundaries
Unknown	Our own "small" unknown overheads	<i>Good luck!</i>

Database client library

- PyMySQL (pure python) vs mariadb (C)
- ~2.5x faster execution just by swapping libraries

	frappe.get_all	frappe.get_doc
Before	1.7	60
After	0.75	22.6

“**`delete`** not possible on large DBs”

1. ``delete`` query on ~95% of data in large tables just fail.
2. No amount of query optimization can fix this.

Fix: Copy 5% data to a temporary table and swap tables.



On *Micro-optimizations*

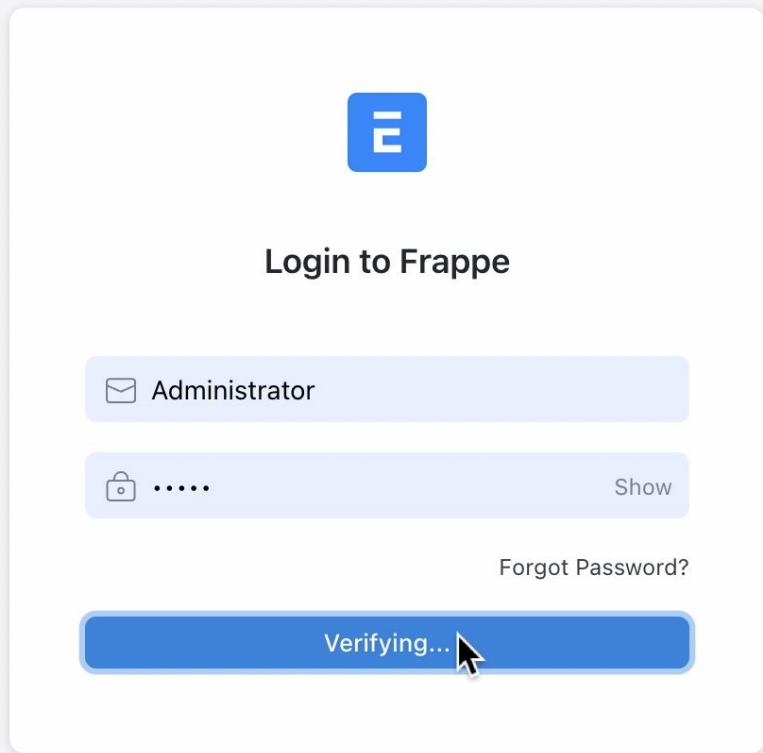
- ``frappe._`` is few microseconds faster.
- ``doc.get`` is few microseconds faster
- Pre-compiling regexps
- ``get_cached_value`` uses dict instead of doc.

“बूँद-बूँद से सागर भरता है”

(The water droplets will one day will an ocean.)

What about feels?

- Numbers don't translate 1-1 with "feels"
- UX improvements:
 - Skeleton loaders
 - Faster splash screens



A mockup of a login screen for 'Frappe'. At the top is a blue square logo with a white 'F'. Below it is the text 'Login to Frappe'. There are two input fields: the first contains 'Administrator' with an envelope icon on the left; the second contains six dots with a lock icon on the left and a 'Show' link on the right. Below the password field is a 'Forgot Password?' link. A large blue button at the bottom says 'Verifying...' with a mouse cursor pointing at it.

Don't have an account? [Sign up](#)

Tools revisited

1. **cProfile** / **%prun** - inbuilt profiler
2. **py-spy** - *sampling* profile best for overheads / in prod.
3. **Browser's Dev tools** - all things client side
4. **memory_profiler** / **scalene** - memory profiling
5. **Frappe Monitor** - monitoring response times and requests
6. **Frappe Recorder** - Analyzing SQL queries in any request.
7. **Custom scripts** - when everything else falls short.

Non-exhaustive list of last year's perf fixes


- 50% faster doc.as_dict
- Faster get_cached_doc()
- Faster workflow actions
- Faster desk routing
- Better DB indexes
- Faster website pages
- Single doc caching
- ...
- Faster BOM update
- Faster reposting
- Faster invoice submission
- Faster PCV submission
- Faster barcode scan
- Faster variant selector
- Faster stock balance report
- ...

Key takeaways

- Trust nothing, profile everything.
- Know thy abstractions.
- Learn some tools, write some tools.

Life isn't easy without them.

- Most performance issues are not deeply *technical*. Just need 🕵️



The Road Ahead

- Performance regression tests
- Persistent DB connections
- Faster Desk routing
- Better client side caching
- More “***scratching the itch***”

★ and watch Frappe for updates.

Questions?